

Class 9: Candy Mini-Project

Alexia (A17297003)

Table of contents

Background	1
Importing Candy Data	1
Exploratory Analysis	4
Overall Candy Rankings	6
Taking a look at pricepercent	12
Exploring the Correlation Structure	14
Principal Component Analysis	15
Summary	17

Background

In this mini-project, we will explore FiveThirtyEight's Halloween Candy data set to find out what is the top ranked snack-sized Halloween candy? and What made some candies more desirable than others? Amongst other questions.

Importing Candy Data

First things first, let's get the data from the FiveThirtyEight GitHub repo. We will download this candy-data.csv file and place it in the project directory. Then, we will load it up with `read.csv()` and inspect the data to see exactly what we're dealing with.

```
candy <- read.csv("candy-data.txt", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

There are 85 different candy types in this data set.

Q2. How many fruity candy types are in the dataset?

```
table(candy$fruity)
```

```
0 1
47 38
```

There are 38 fruity candy types in the data set.

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

```
candy["Snickers", ]$winpercent
```

```
[1] 76.67378
```

Q4. What is the winpercent value for "Kit Kat"?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

Q5. What is the `winpercent` value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

Side Note: the `skimr::skim()` function There is a useful `skim()` function in the `skimr` package that can help give us a quick overview of a given data set. We will install this package and try it on our candy data.

```
library("skimr")
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyal- mondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice- wafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	

skim_vari- able	n_miss- ing	com- plete_rate	mean	sd	p0	p25	p50	p75	p100	hist
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

The following columns look to be on a different scale to the majority of the other columns in the data set: sugarpercent, pricepercent, and winpercent.

Q7. What do you think a zero and one represent for the candy\$chocolate column?


I think the zero and one represent how much of the data is missing and how much of the data is measured, respectively.

Exploratory Analysis

A good place to start any exploratory analysis is with a histogram. You can do this most easily with the base R function `hist()`. Alternatively, you can use `ggplot()` with `geom_hist()`

Q8. Plot a histogram of `winpercent` values using both base R and ggplot2.

```
hist(candy$winpercent)
```



Q9. Is the distribution of `winpercent` values symmetrical?

The distribution of `winpercent` values is not symmetrical.

Q10. Is the center of the distribution above or below 50%?

```
median(candy$winpercent)
```

```
[1] 47.82975
```

The center of distribution is below 50%.

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
choc.ind <- as.logical(candy$chocolate)
choc.candy <- candy[choc.ind, ]
choc.win <- choc.candy$winpercent
mean(choc.win)
```

```
[1] 60.92153
```

```
fruity.ind <- as.logical(candy$fruity)
fruity.candy <- candy[fruity.ind, ]
fruity.win <- fruity.candy$winpercent
mean(fruity.win)
```

```
[1] 44.11974
```

On average chocolate candy is ranked higher than fruit candy.

Q12. Is this difference statistically significant?

```
t.test(choc.win, fruity.win)
```

Welch Two Sample t-test

```
data:  choc.win and fruity.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

This difference is statistically significant.

Overall Candy Rankings

Let's use the base R `order()` function together with `head()` to sort the whole data set by a variable/column of interest.

Q13. What are the five least liked candy types in this set?

```
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip				0	0	0	1	0.197		0.976
Boston Baked Beans				0	0	0	1	0.313		0.511
Chiclets				0	0	0	1	0.046		0.325
Super Bubble				0	0	0	0	0.162		0.116
Jawbusters				0	1	0	1	0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

The 5 least liked candies in this data set are Nik L Nip, Boston Baked Beans, Chiclets, Super Bubble and Jawbusters.

Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy[order(-candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Reese's Peanut Butter cup	1	0	0		1	0
Reese's Miniatures	1	0	0		1	0
Twix	1	0	1		0	0
Kit Kat	1	0	0		0	0
Snickers	1	0	1		1	1

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent
Reese's Peanut Butter cup				0	0	0		0.720
Reese's Miniatures				0	0	0		0.034
Twix				1	0	1	0	0.546
Kit Kat				1	0	1	0	0.313
Snickers				0	0	1	0	0.546

	price	percent	winpercent
Reese's Peanut Butter cup	0.651		84.18029
Reese's Miniatures	0.279		81.86626
Twix	0.906		81.64291


Kit Kat	0.511	76.76860
Snickers	0.651	76.67378

The 5 most liked candies in the data set are Reese's Peanut Butter cup, Reese's Miniatures, Twix, Kit Kat, and Snickers.

Q15. Make a first barplot of candy ranking based on winpercent values.


```
library(ggplot2)

ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```



Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?


```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```

Let's setup a color vector (that signifies candy type) that we can then use for some future plots.

```
mycols <- rep("black", nrow(candy))
mycols[as.logical(candy$chocolate)]="chocolate"
mycols[as.logical(candy$bar)]="brown"
mycols[as.logical(candy$fruity)]="pink"
```

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy),winpercent)) +
  geom_col(fill=mycols) +
  ylab("")
```



Q17. What is the worst ranked chocolate candy?

The worst ranked chocolate candy is Sixlets.

Q18. What is the best ranked fruity candy?


The best ranked fruity candy is Starburst.

Taking a look at pricepercent

What about value for money? What is the best candy for the least money? One way to get at this would be to make a plot of `winpercent` vs the `pricepercent` variable.

```
library(ggrepel)

ggplot(candy) +
  aes(winpercent,
      pricepercent,
      label=rownames(candy)) +
  geom_point(col=mycols) +
  geom_text_repel(col=mycols, size = 3.3, max.overlaps=5)
```



Q19. What about value for money? What is the best candy for the least money?

```
head(candy[order(candy$pricepercent),], n=3)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
Tootsie Roll Midgies	1	0	0	0	0
Pixie Sticks	0	0	0	0	0
Dum Dums	0	1	0	0	0

	crispedricewafer	hard bar	pluribus	sugarpercent
Tootsie Roll Midgies	0	0	1	0.174
Pixie Sticks	0	0	1	0.093
Dum Dums	0	1	0	0.732

	pricepercent	winpercent
Tootsie Roll Midgies	0.011	45.73675
Pixie Sticks	0.023	37.72234
Dum Dums	0.034	39.46056

The best candy for the least amount of money is Tootsie Roll Midgies.

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
head(candy[order(-candy$pricepercent),], n=5)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat
Nik L Nip	0	1	0	0	0
Nestle Smarties	1	0	0	0	0
Ring pop	0	1	0	0	0
Hershey's Krackel	1	0	0	0	0
Hershey's Milk Chocolate	1	0	0	0	0

	crispedricewafer	hard bar	pluribus	sugarpercent
Nik L Nip	0	0	1	0.197
Nestle Smarties	0	0	1	0.267
Ring pop	0	1	0	0.732
Hershey's Krackel	1	0	1	0.430
Hershey's Milk Chocolate	0	0	1	0.430

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

The top 5 most expensive candy types are Nestle Smarties, Ring Pop, Nik L Nip, Hershey's Krackel, and Hershey's Milk Chocolate. The least popular amongst all of these is Nik L Nip.


Exploring the Correlation Structure

Before we dive into PCA in the next section, it's worth pausing to look more closely at the correlations between all pairs of variables.

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
cij <- cor(candy)  
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

The two variables that are anti-correlated are fruity and chocolate, fruity and bar, pluribus and bar.

Q23. Use your corrplot result to make a prediction: which variables do you expect will have the largest contributions (i.e. loadings) to PC1 (i.e., drive the most separation between candies along the first principal component)?

The variables we could expect to have the largest contributions to PC1 are fruity and chocolate.

Principal Component Analysis

In this case we want to be sure to set `scale = TRUE` argument for `prcomp()` because we have one variable `winpercent` that is on a very different scale than all others and would otherwise dominate our PCA results

```
pca <- prcomp(candy, scale = TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369


	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

First major result figure is the “score plot” of PC1 vs PC2 - how the different candy are related to each other on our new PC axis:

```
ggplot(pca$x) +
  aes(PC1, PC2, label=row.names(pca$x))+
  geom_point(col=mycols)+
  geom_text_repel(col=mycols)
```



```
ggplot(pca$rotation) +
  aes(PC1,
       reorder(row.names(pca$rotation), PC1)) +
  geom_col()
```

Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

The original variables that are picked up strongly by PC1 in the positive direction are fruity, pluribus, and hard. This makes sense as we saw in the corrplot that there is no correlation between fruity and chocolate therefore the divide in the data would be significant.

Summary

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

The following combination of characteristics appears to make a “winning” candy: chocolate, bar, peanuts, and nougat.

The plots generated with the data help identify these findings as they show us multiple times that there is no correlation between fruity and chocolate (corrplot) and also visually demonstrate that the most popular candy type between chocolate and fruity is in fact chocolate (Visualization).